
ISA 414 – Managing Big Data

Lecture 12 – Data Analysis

Supervised Learning (Part I)

Dr. Arthur Carvalho

arthur.carvalho@miamioh.edu



MIAMI UNIVERSITY

Copyright © 2021 Arthur Carvalho

Announcements

➤ Assignment 3

- Now available on Canvas
- Deadline: Wednesday, 10/06, before 11:59 pm

Lecture Objectives

- Quick review of Homework 5 and 6
- Learn about predictive analytics
 - Classification and regression problems
 - Decision trees
 - Random forests
 - How to evaluate models
 - Training and test sets

Lecture Instructions

- Download the following files available on Canvas
 - *Lecture 12.ipynb*
 - *energy_data.csv*
 - *Description of the variables.xlsx*
- Place the above files inside the same folder
- Open the file *Lecture 12.ipynb* with VS Code

CRISP-DM



➤ Business Understanding phase

- Important question: can a certain business problem be formulated as a data-analytics problem?

➤ Different types of problems/answers

- Descriptive analytics: provides answers on “*what happened*”
 - Descriptive statistics, SQL queries, OLAP (business intelligence), ...
- Prescriptive analytics: provide answers on “*what could have happened*”
 - Optimization, simulation, ...
- Predictive analytics: provide answers on “*what will happen*”
 - Focus of this lecture

Business Understanding

- The solutions to the category of problems we will be working on are based on a paradigm called supervised learning
 - Assumption: there is a clear target variable whose values we are trying to predict

	Problem	Supervised	Unsupervised
Our focus today →	Classification	X	
Next class →	Regression	X	
	Clustering		X
	Co-occurrence grouping		X
	Profiling		X

Supervised Learning

➤ Classification (or class probability estimation)

- Goal: determine which class a new observation (likely) belongs to
 - The target variable is qualitative (categorical/factor)
- Example (Assignment 3): given a data set containing data about bank clients, can we build a statistical model that classifies future clients as “good” or “bad” in terms of paying back a loan?
 - Clear categorical target (two values: “good” and “bad”)

Supervised Learning

➤ Regression problems

- Goal: predict the value of a target variable using a model structure that relates the target with informative attributes
 - The target variable is quantitative (numeric)
- Example: what is the market price of a house in Oxford with 4 bedrooms, 3 bathrooms, built in 2016, ...
 - Clear numeric target (price)

Supervised Learning

➤ Illustrative problem

- An energy (electricity) provider has historical data with information about current clients
 - *E.g.*, Age, salary, marital status, *etc.*
- Each client (observation) in the data set is associated with a category
 - Electricity tariff (*e.g.*, time-of-use tariff, peak tariff, flat tariff)
- Business question:
 - Which electricity tariff should this company suggest to a new client?
 - Is this a classification or a regression problem?


Classification

➤ Potential solution:

- Translate this problem into a classification problem
- Collect the data
 - *E.g.*, query internal databases
- Build a *classifier*
 - A machine learning model that classifies clients based on tariffs
- Classify the new client
 - Using the classifier to estimate which category the new client likely belongs to

Classification


- Consider the following hypothetical situation:
 - All clients younger than x in your data set prefer Tariff 1
 - All clients older than x in your data set prefer Tariff 2
- What would be a good classifier?
 - Which tariff would you suggest to a new client?



Classification


➤ Classifier:

- New client's age $> x$
 - Suggest Tariff 2
- Otherwise
 - Suggest Tariff 1



Classification


- What about this more realistic case?
 - Two variables



Classification


➤ Classifier:

- New client's age $> x$ AND new client's salary $> y$
 - Suggest Tariff 2
- New client's age $> x$ AND new client's salary $< y$
 - Suggest Tariff 1
- New client's age $< x$ AND new client's salary $> y$
 - Suggest Tariff 1
- New client's age $< x$ AND new client's salary $< y$
 - Suggest Tariff 2



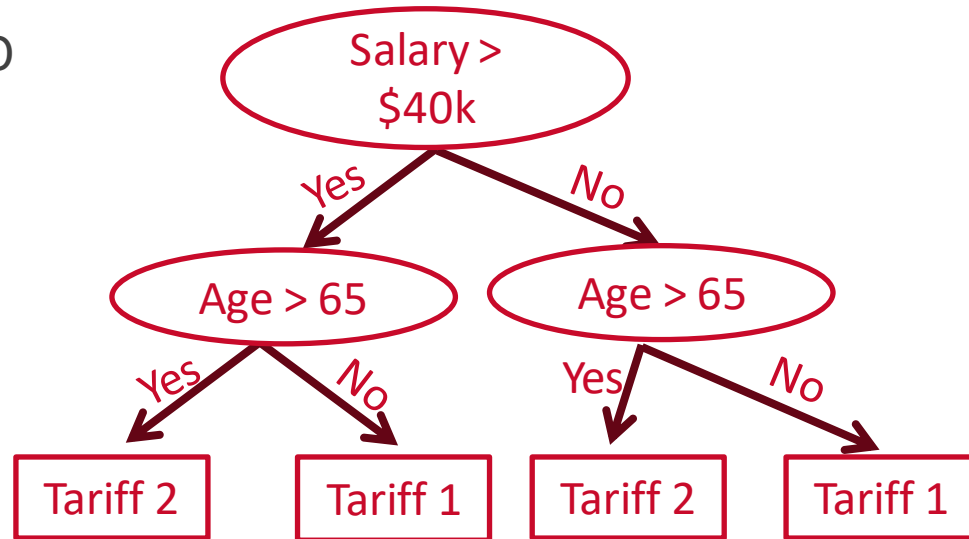
Classification

- Keep in mind that real-life is a mess



Decision Trees

- The previous approach of segmenting the attribute space is precisely what a technique called *Decision Trees* does
 - Create “if-then” rules
E.g., IF ‘Salary > \$40K’ AND
‘Age > 65’
THEN suggest Tariff 2



Decision Trees

- Formally, decision trees create hyperrectangles
- Categories are the “leaves”
 - Bottom of the tree
- There are many algorithms for building decision trees from a data set
 - Beyond the scope of this course
 - General idea: choose a variable at each level that best splits the data set
 - *E.g.*, reduce entropy (information gain)

Decision Trees

- Henceforth, we shall heavily use the `sklearn` and `pandas` modules in our modeling efforts
 - Note that `sklearn` has several submodules that serve different purposes
 - We shall import them as we progress
 - Install the required modules

```
pip install sklearn
pip install pandas
```
 - Load `pandas` and the data set we use today

```
import pandas as pd
energy_data = pd.read_csv("energy_data.csv")
```

Decision Trees

- Before analyzing the data, always check whether each attribute is of the right type
 - If it is not, then change attribute types
 - Note that pandas store strings as objects

`energy_data.dtypes`
- **sklearn** does not allow for non-numeric variables when building models
 - Its `preprocessing` submodule has several functions to transform categorical into continuous variables
 - Example: `OrdinalEncoder()`, `OneHotEncoder()`, `LabelEncoder()`

Decision Trees

- Oftentimes (not always) one should encode categorical variables as **dummies**
 - Except for the **target** variable, which should have its values replaced by numbers
- Let's derive dummies for the predictors
 - `pandas` offer a simpler way to derive dummies than `sklearn`

```
energy_data = pd.get_dummies(energy_data,  
                               columns = ["MaritalStatus", "IncomeLevel", "DwellingArea",  
                                          "HasChildren", "SolarRoof", "ShiftableLoad",  
                                          "AttitudeSustainability" ],  
                               drop_first = True)
```

Decision Trees

- Let's derive recode the target variable
 - (Sub)module `preprocessing` in `sklearn`

```
from sklearn import preprocessing
```

```
enc = preprocessing.LabelEncoder()
```

```
energy_data["Tariff"] = enc.fit_transform(energy_data["Tariff"])
```

Decision Trees

➤ Building the tree model

- One must create two sets of columns

- The feature (independent) variables

```
x = energy_data.drop(columns=["Tariff"])
```

- The target (dependent) variable

```
y = energy_data["Tariff"]
```

- Next, it is time to create and fit a model

```
model = DecisionTreeClassifier()
```

```
model = model.fit(x,y)
```

Model Evaluation

- How do we know our model is any good?
 - One possible way: split the original data set into two parts
 - Train the model using one data set (training set)
 - Test the model using the complementary data set (test set)
 - Use the trained model to predict the target values in the test set, and compare the predictions against the true values
 - The higher the number of times the predictions agree with the true values, the more accurate the classifier is

Model Evaluation

➤ Analogy: exam

- A professor gives you a study guide
 - Set of problems with answers
- What if the professor asks you the same questions in the exam?
 - The professor is not really testing your knowledge
 - The professor is testing whether you are capable of memorizing answers
- What if the professor asks you slightly different questions about the same material in the exam?
 - The professor is now testing your knowledge (generalization power)
- Study guide = training; Exam = testing

Model Evaluation

- Let's redo what we did before and evaluate our model
 - Randomly split the original data into two data frames
 - **Training set** (66% of the observations)
 - **Test set** (34% of the observations)
 - The 66/34 division is just one common way of doing it
 - Train the model using the training set
 - Evaluate the model using the test set

Model Evaluation

- Let's use the submodule `model_selection` in `sklearn` to split a data frame

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.34)
```

- Building a decision tree using the training set

```
model = DecisionTreeClassifier()
```

```
model = model.fit(x_train,y_train)
```

Model Evaluation

- Evaluating a decision tree using the test set
 - The `metrics` submodule contains several metrics to evaluate statistical models
 - We shall use the overall accuracy metric
 - Percentage of correctly classified instances
 - `from sklearn import metrics`
 - Step 1: use the model to predict the class of each observation in the test set
`y_pred = model.predict(x_test)`
 - Step 2: calculate how often the predictions in the model agree with the true class in the test set
`metrics.accuracy_score(y_test, y_pred)`
 - An accuracy of, say, 0.60 means that the model is expected to correctly classify 60% of future instances

Decision Trees

➤ Strengths

- Simple to understand and interpret
- Performs reasonably well for big data sets

➤ Drawbacks

- Learning the optimal tree is not always computationally feasible
- Might result in a high bias towards the training set
 - Tendency to **overfit**

Classification

- There are many models for classification
 - SVM, random forests, GBM, logistic regression...
 - Which one is the best?
 - In theory, all algorithms are equally good in expectation !!!
 - No Free-Lunch Theorem
 - Common approach when predictive accuracy is the only important factor
 - Build and evaluate multiple classifiers
 - Perform statistical analysis on the obtained results to determine the most accurate model


Classification

- Some models perform well for certain problems, and poorly for other problems
- Another common approach: combine several models
 - Ensemble learning
 - Compensate poor individual performance
 - (Almost) free lunch
 - Diversity matters
 - Formally, one wants the errors produced by the individual models to be as uncorrelated as possible


Random Forests

- Ensemble model
- Idea: build several decision trees semi-randomly
 - Each tree individually classifies a new observation
 - The most popular predicted category is chosen as the outcome of the model


Tariff 1



Tariff 2



Tariff 1



Random Forests

- Building a random forest (200 trees) using the training set

- Submodule `sklearn.ensemble`

```
from sklearn.ensemble import RandomForestClassifier  
model = RandomForestClassifier(n_estimators=200)  
model = model.fit(x_train,y_train)
```

- Evaluating a random forest model using the test set

```
y_pred = model.predict(x_test)  
metrics.accuracy_score(y_test,y_pred)
```


Random Forests

- One of the most popular models in forecasting competitions (alongside GBM and Neural Networks)
 - Knowledge Discovery in Databases (KDD)
 - Kaggle.com
- Strengths
 - Tackles the bias problem with single decision trees
- Drawbacks
 - No longer easy to interpret and explain the results

Classification

- This lecture summarized what is often taught across many data mining classes
- Keep in mind that:
 - There are many different statistical models for different types of problems
 - There are many different evaluation metrics other than using the percentage of correctly classified observations
 - *E.g.*, specificity, sensitivity, ROC area
 - There are many different ways of estimating model errors
 - K-fold cross validation, nested cross validation

Summary

➤ Summary

- Data-analytics problems: classification problem
- Decision trees and random forests
- Evaluation: training and test sets

➤ Useful references

- <http://www.r2d3.us/visual-intro-to-machine-learning-part-1/>
- <https://docs.google.com/presentation/d/1kSuQyW5DTnkVaZEjGYCkfOxvzCqGEFzWBy4e9Ueddd9k/edit>

➤ Next class: regression problems

Copyright 2021 Arthur Carvalho. All rights reserved. This material may not be published, broadcast, rewritten, or redistributed without explicit written consent.