
ISA 414 – Managing Big Data

Lecture 22 – Introduction to Hadoop

(Part II)

Dr. Arthur Carvalho

arthur.carvalho@miamioh.edu

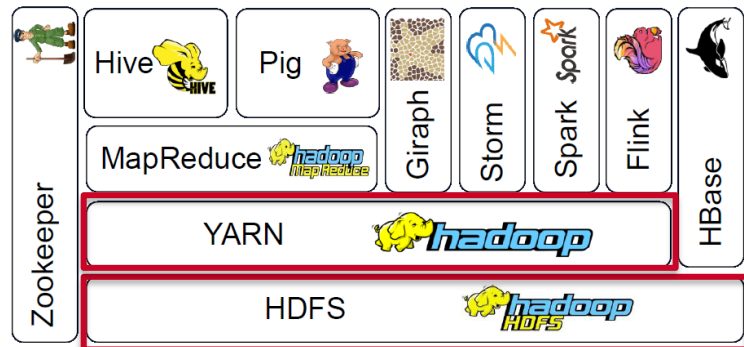


MIAMI UNIVERSITY

Copyright © Arthur Carvalho

Lecture Objectives

- Understand the inner workings of Hadoop
 - HDFS



Lecture Instructions

- There are several new concepts today
 - Suggestion: actively take notes
 - Important keywords are highlighted in the slides

Hadoop

- Summary of previous lecture
 - Handling massive amounts of data is one of the biggest challenges brought by big data
 - How to store massive amounts of data?
 - Cheap solution: **commodity clusters** and **distributed storage**
 - Individual files are stored in different computers in a cluster
 - Massive files are broken down into chunks of data, which are stored in different computers (nodes) in a cluster
 - A **distributed file system** tracks where each piece of data is located

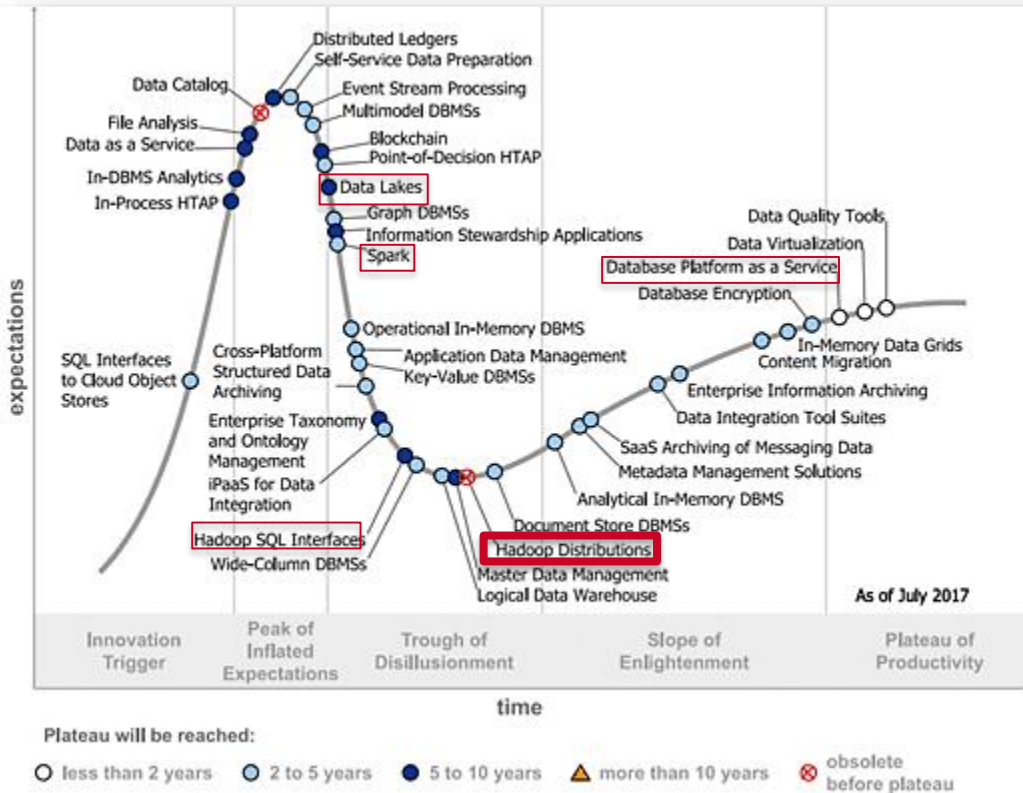
Hadoop

- Summary of previous lecture
 - Handling massive amounts of data is one of the biggest challenges brought by big data
 - How to analyze (perform computations using) big data?
 - Cheap solution: **commodity clusters** and **distributed computation**
 - Paradigm: “**move computation to data**”

Hadoop

- Summary of previous lecture
 - **Hadoop:** framework used for distributed storage and computing
 - */i.e.*, a tool that manages commodity clusters
 - A collection of technologies
 - From data storage to data collection and analysis, these technologies might have drastically different roles
 - Many of the Hadoop technologies are still part of Gartner's Hype Cycle
 - https://blogs.gartner.com/andrew_white/2020/08/07/data-and-analytics-hype-cycles-for-2020-just-published/

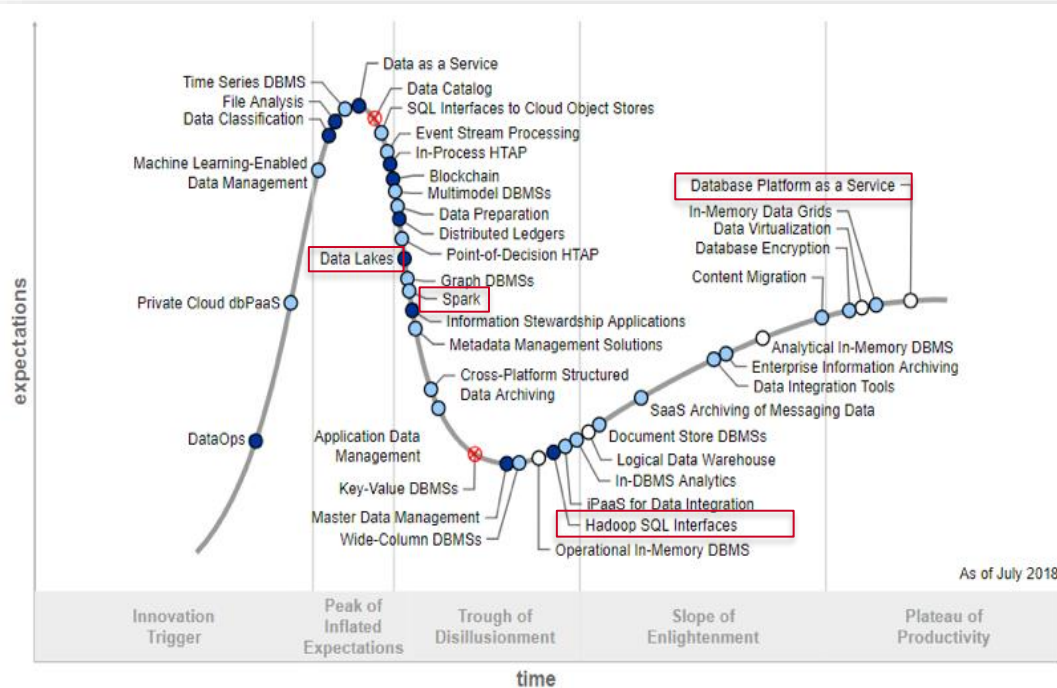
Hadoop



Gartner 2017: Hype Cycle for Data Management

(source: <https://www.gartner.com/newsroom/id/3809163>)

Hadoop



Gartner 2018: Hype Cycle for Data Management

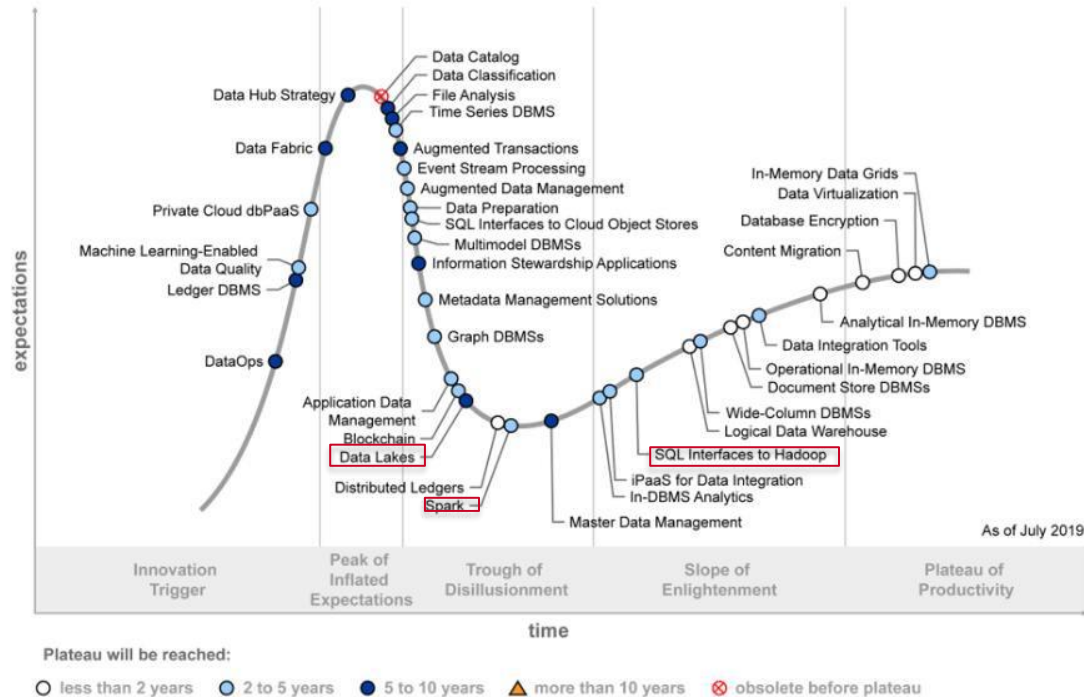
(source: <https://www.gartner.com/en/newsroom/press-releases/2018-09-11-gartner-hype-cycle-for-data-management-positions-three-technologies-in-the-innovation-trigger-phase-in-2018>)

Plateau will be reached:

○ less than 2 years ● 2 to 5 years ● 5 to 10 years ▲ more than 10 years ✕ obsolete before plateau

Hadoop

Hype Cycle for Data Management, 2019



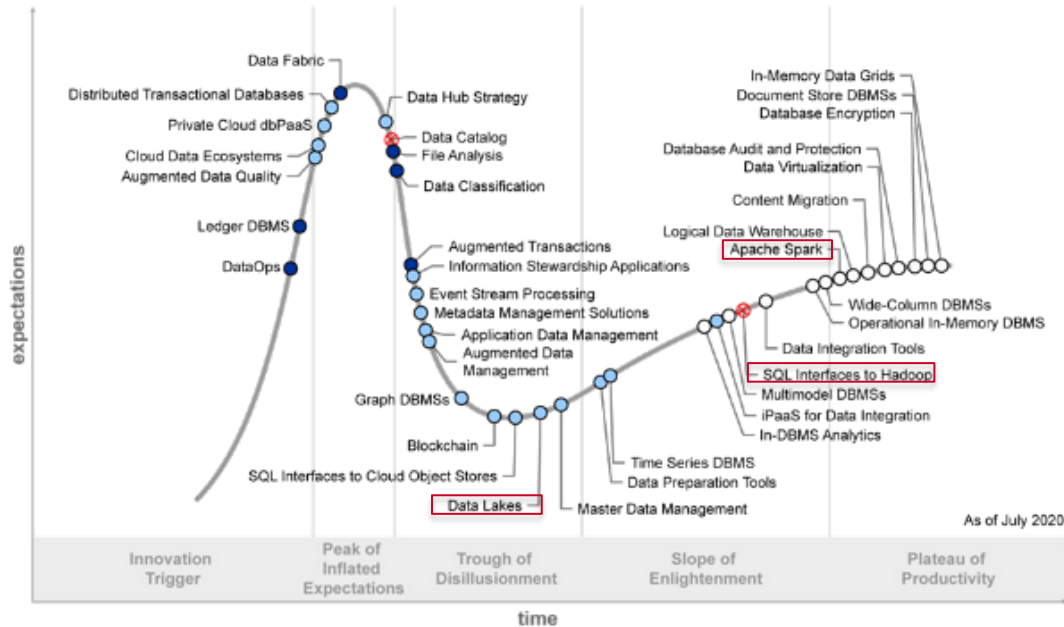
Gartner 2019: Hype Cycle for Data Management

(source:

<https://www.gartner.com/en/documents/3955768/hype-cycle-for-data-management-2019>)

Hadoop

Hype Cycle for Data Management, 2020



Gartner 2020: Hype Cycle for Data Management

(source: <https://www.denodo.com/en/document/analyst-report/gartner-hype-cycle-2020>)

HDFS

➤ Hadoop Distributed File System (HDFS)


- Remember: **distributed storage** requires a **distributed file system**
 - That is what **HDFS** is
- Main idea: HDFS breaks down large files into file blocks and spread them across multiple computers in a cluster
- Storage layer
 - Foundation for most tools in the Hadoop ecosystem
 - Scalable: one can easily add more nodes to increase total storage space
 - Reliable: fault tolerant (more on this soon)

HDFS

- Think about a very large data file (gigabytes or petabytes)
 - Example: Kaggle's Data Science Bowl 2017
 - Goal: improve lung cancer detection
 - Prize: \$1,000,000
 - Data size: 67+ GB
- HDFS breaks such files into chunks (blocks) and spread them over a computer cluster

HDFS


- Example: suppose a person uploads a 256-MB file to HDFS
 - Default block (piece) size: 64 MB
 - Configurable
 - That file is split into four parts ($256/64 = 4$) and the resulting blocks are spread across the cluster



- In the above example, what happens if Node C fails?
 - *E.g.*, power outage

HDFS

- HDFS is designed for fault tolerance
 - Replication: HDFS makes copies of blocks on different nodes to prevent data loss
 - Default: 3 copies (configurable)



Example: if a node containing block C fails, two other nodes can still provide block C for a requesting application


HDFS

➤ HDFS vs Relational Databases Vs MongoDB

- HDFS stores files (virtually any type)
 - Structured (e.g., CSV) and unstructured (e.g., images) data are inside files
- Relational databases store data inside tables
 - Highly structured data
 - One must design a database model first
 - Not very flexible
- MongoDB stores data using the JSON format
 - Very flexible when handling textual data
 - No need for a predefined model

HDFS

➤ Technical aspects: bird's-eye view



Data Nodes stores the data and return it on request


HDFS

- Technical aspects: two key components
 - **Name Node** for metadata
 - Handles metadata (data about the cluster)
 - Manager of the HDFS cluster
 - Keeps track of file names, location of blocks in directories/computers, *etc.*
 - Usually one per cluster, but there might exist a secondary, backup node
 - **Data Node** for block storage
 - Nodes store data (file blocks)
 - Usually, one Data Node refers to one machine
 - So, there are often many Data Nodes per cluster
 - Listens to commands from the Name Node
 - Block creation, deletion, replication

HDFS

➤ Simplified example: suppose a task (client) wants to write a big file to a Hadoop cluster

1. Client makes a request to store Block 1;
Name node informs the client the locations (e.g., IP addresses) where Block 1 must be stored
 - First location: Data Node 1 (DN 1)
2. Client sends Block 1 and a list of locations to DN 1
3. DN 1 stores Block 1 and sends Block 1 plus a list of locations to DN 2
4. DN 2 stores Block 1 and sends Block 1 plus a list of locations to DN 4



HDFS

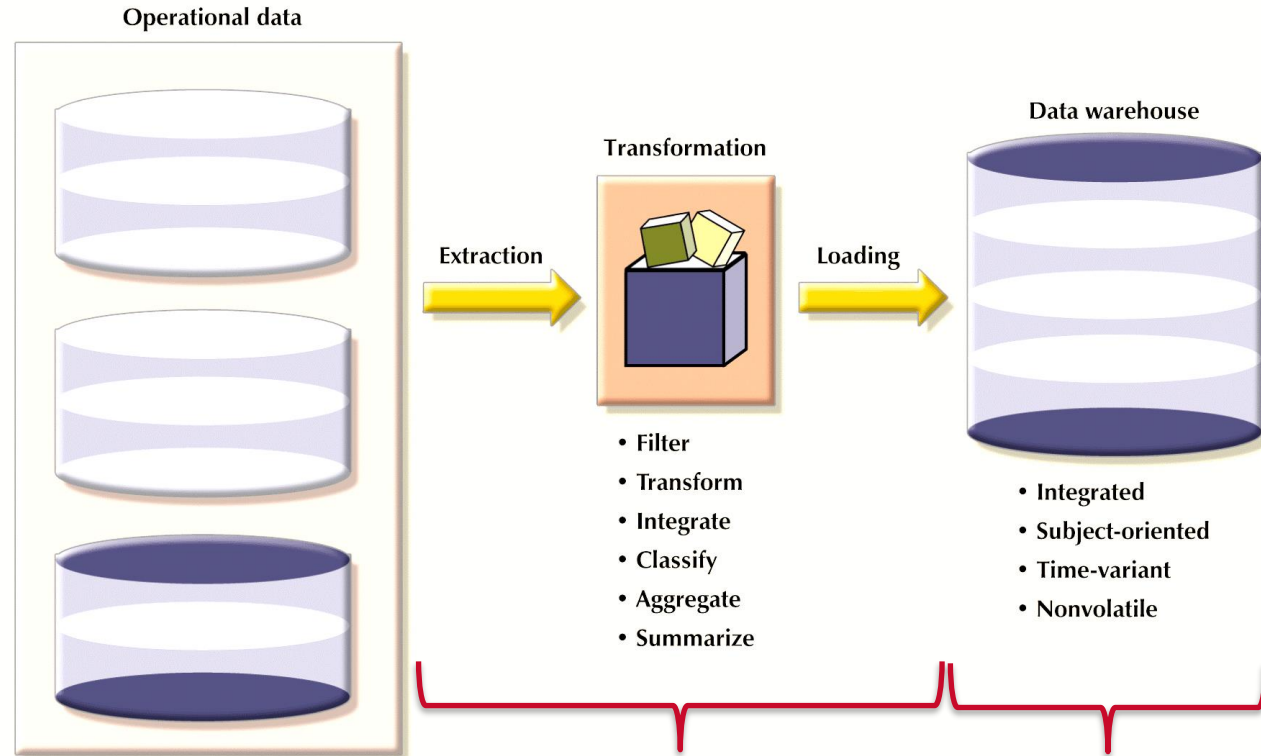
- HDFS is all about storing files in a distributed manner
 - We will soon observe how HDFS works in practice
 - Before doing do, let's review some relevant concepts
 - Data warehouse
 - Data lake

Data Warehouse

➤ Data Warehouse

- Business intelligence (BI) technique to integrate and analyze data
- Oftentimes, it is about an enormous collection of data
 - *Subject oriented* – designed around key entities (concepts)
 - *Integrated* – consistency in naming convention, encoding, translation, ...
 - *Time-variant* – data are organized by time periods
 - *Non-volatile* – data are updated in batches, rather than as transactions occur

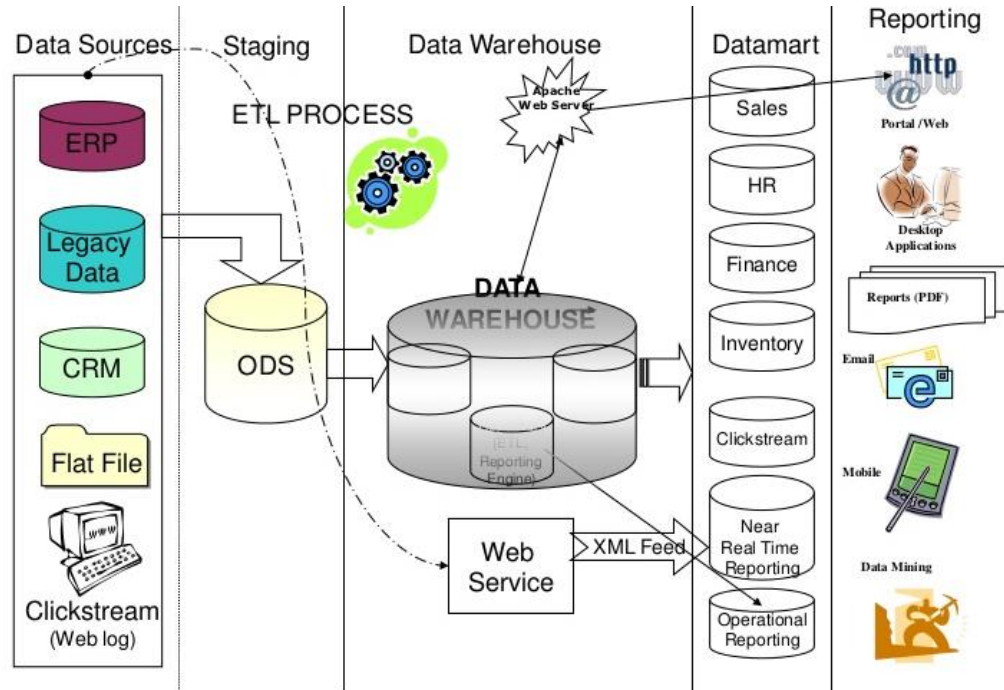
Data Warehouse



Source:
Database Systems,
9th Edition

Data Warehouse

➤ A more realistic view



Data Warehouse

➤ Operational data

- Mostly stored in relational databases
- Optimized to support transactions representing daily operations

➤ Decision support data

- Derived from operational data
- Stored in a database optimized for data analysis and query processing

Data Warehouse

- **ETL** - key processes for deriving decision support data from operational data
 - **EXTRACT**
 - The process of extracting operational data, e.g., query a relational database
 - Can be very time-consuming
 - Often happens when the operational databases are not heavily used
 - Data can be stored in distributed databases
 - Potentially different database vendors (e.g., Oracle, MySQL) or representations (e.g., relational model, network model, hierarchical model)

Data Warehouse

➤ **ETL** - key processes of deriving decision support data from operational data

- **TRANSFORM**

- Standardizing different pieces of data
 - Example: convert “male”/“female” to 0/1, inches to centimeters, date format from “dd/mm/yy” to “yyyy/mm/dd”
- Cleaning data
 - Same record stored in different databases
 - Misuse of data entry fields
 - *E.g.*, age = 200, SSN = 0000000000

Data Warehouse

- **ETL** - key processes of deriving decision support data from operational data
 - **LOAD**
 - The process of loading the data into the warehouse
 - Occur in batches, rather than after each transaction involving operational data

Data Warehouse

- Is HDFS a data warehouse technology?
 - Technically, no!

- Unlike HDFS, a data warehouse:
 1. Follows a pre-defined architecture
 - *E.g.*, a star schema having dimensions, facts, and attributes
 2. Follows a strong data quality assurance process (*i.e.*, ETL)

Data Lake

➤ So, what is HDFS?

- Answer: a ***Data Lake***

- Centralized repository of data
- Stores all kinds of data in raw format (*i.e.*, files)
 - Images, texts, audio, ...
- Data are not necessarily curated before being dumped into a data lake
 - One of the main criticism of Hadoop HDFS

Data Lake


➤ So, what is Hadoop HDFS?

- Answer: a *Data Lake*
 - Hadoop was “the next big thing” about 10 years ago
 - Common approach followed by many companies during Hadoop’s hype peak time:

“We see customers creating big data graveyards, dumping everything into HDFS [Hadoop Distributed File System] and hoping to do something with it down the road. But then they just lose track of what’s there. The main challenge is not creating a data lake, but taking advantage of the opportunities it presents” (Sean Martin, CTO of Cambridge Semantics)

HDFS

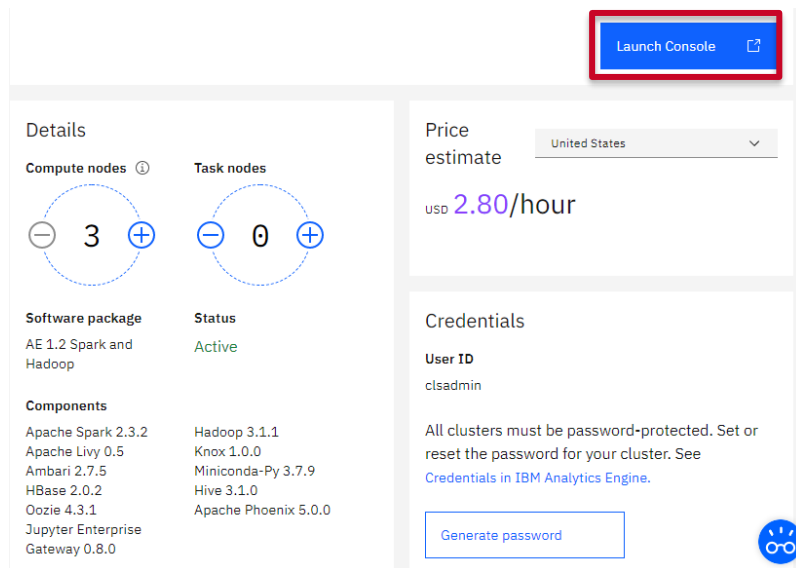
- Let's visualize some of the previous concepts
 - Recall that HDFS is a back-end technology
 - “Boring,” not visually appealing
 - Cluster created on IBM Cloud
 - PaaS called *Analytics Engine*
 - See our previous class for instructions on how to set up this service
 - 1 Name Node (4 cores), 3 Data Nodes (12 cores)
 - \$2.8 per hour
 - Topology



HDFS

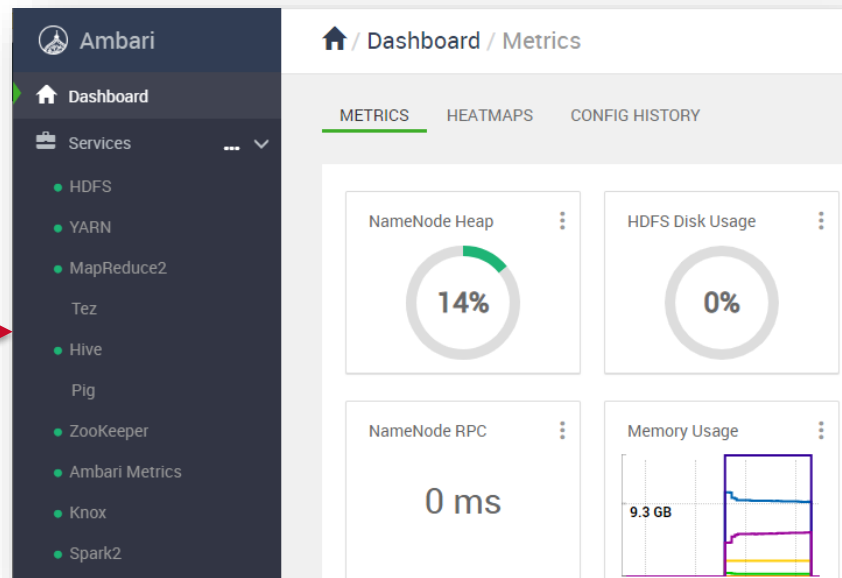
➤ Cluster administrator's perspective

- Ambari software



The image shows the Ambari console configuration page for a new cluster. A red box highlights the "Launch Console" button in the top right corner. The page is divided into several sections:

- Details:** Includes "Compute nodes" (3) and "Task nodes" (0) with minus and plus buttons. It also shows the "Software package" (AE 1.2 Spark and Hadoop) and "Status" (Active).
- Components:** Lists installed components: Apache Spark 2.3.2, Apache Livy 0.5, Ambari 2.7.5, HBase 2.0.2, Oozie 4.3.1, Jupyter Enterprise, and Gateway 0.8.0.
- Price estimate:** Shows a price of USD 2.80/hour for the United States.
- Credentials:** Displays the "User ID" as clsadmin and provides a link to "Generate password".



The image shows the Ambari dashboard metrics page. The left sidebar lists services: HDFS, YARN, MapReduce2, Tez, Hive, Pig, ZooKeeper, Ambari Metrics, Knox, and Spark2. The main content area displays various metrics:

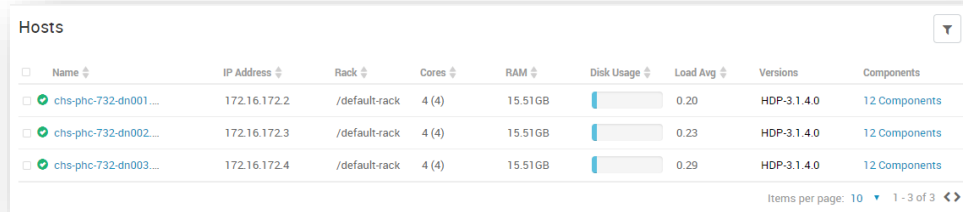
- NameNode Heap:** 14% (represented by a circular gauge).
- HDFS Disk Usage:** 0% (represented by a circular gauge).
- NameNode RPC:** 0 ms.
- Memory Usage:** 9.3 GB (represented by a line graph).

HDFS

➤ Cluster administrator's perspective

- Ambari software

- Services: Hadoop services running in the cluster
- *Services -> HDFS -> DATANODES*: Info about data nodes




The screenshot shows the 'Hosts' page in the Ambari web interface. It displays a table with columns for Name, IP Address, Rack, Cores, RAM, Disk Usage, Load Avg, Versions, and Components. Three data nodes are listed, each with a green status icon and a link to '12 Components'.

<input type="checkbox"/>	Name	IP Address	Rack	Cores	RAM	Disk Usage	Load Avg	Versions	Components
<input checked="" type="checkbox"/>	chs-phc-732-dn001...	172.16.172.2	/default-rack	4 (4)	15.51GB	<div><div></div></div>	0.20	HDP-3.1.4.0	12 Components
<input checked="" type="checkbox"/>	chs-phc-732-dn002...	172.16.172.3	/default-rack	4 (4)	15.51GB	<div><div></div></div>	0.23	HDP-3.1.4.0	12 Components
<input checked="" type="checkbox"/>	chs-phc-732-dn003...	172.16.172.4	/default-rack	4 (4)	15.51GB	<div><div></div></div>	0.29	HDP-3.1.4.0	12 Components

Items per page: 10 | 1 - 3 of 3

- *Views -> "Files Views"*: Upload data to HDFS



HDFS

- There are 2 files inside my cluster
 - `/user/clsadmin/Data.csv` (1.2 MB)
 - `/user/clsadmin/ Docklands.jmp` (450 MB)
- What happened behind the scenes when I uploaded these files?
 - Tip: IBM uses blocks of size 128 MB; replication factor = 3
- Hacking time
 - I will connect to the Name Node using SSH
 - Issue two commands to get info about the above files

```
hdfs fsck /user/clsadmin/Data.csv -files -locations -blocks
```

```
hdfs fsck /user/clsadmin/Docklands.jmp -files -locations -blocks
```

HDFS

File #1: Data.csv (1.2 MB)

FSCK started by clsadmin (auth:SIMPLE) from /172.16.122.135 for path /user/clsadmin/Data.csv at Tue Mar 09 15:48:52 UTC 2021
/user/clsadmin/Data.csv 1230849 bytes, replicated: replication=3, 1 block(s): OK
0. BP-2089008039-172.16.122.134-1615302408196:blk_1073743602 2778 len=1230849 live_repl=3 [DatanodeInfoWithStorage[172.16.122.131:50010,DS-f4d957e2-1c71-42e9-8310-c5d0aec12cb7,DISK], DatanodeInfoWithStorage[172.16.122.132:50010,DS-df4aac5a-b829-4b47-adbe-b5d79cbc75de,DISK], DatanodeInfoWithStorage[172.16.122.130:50010,DS-126c24f2-174c-46ca-b9bb-7ba343d48a2c,DISK]]

Status: HEALTHY

Number of data-nodes: 3
Number of racks: 1
Total dirs: 0
Total symlinks: 0

Replicated Blocks:

Total size: 1230849 B
Total files: 1
Total blocks (validated): 1 (avg. block size 1230849 B)
Minimally replicated blocks: 1 (100.0 %)
Over-replicated blocks: 0 (0.0 %)
Under-replicated blocks: 0 (0.0 %)
Mis-replicated blocks: 0 (0.0 %)
Default replication factor: 3
Average block replication: 3.0
Missing blocks: 0
Corrupt blocks: 0
Missing replicas: 0 (0.0 %)

HDFS

File #2: Docklands.jmp (450 MB)

```
FSCCK started by clisadmin (auth:SIMPLE) from /172.16.122.135 for path /user/clisadmin/Docklands.jmp at Tue Mar 09 15:52:12 UTC 2021
/user/clisadmin/Docklands.jmp 450200813 bytes, replicated: replication=3, 4 block(s): OK
0. BP-2089008039-172.16.122.134-1615302408196:blk_1073743598_2774 len=134217728 Live_repl=3 [DatanodeInfoWithStorage[172.16.122.130:50010,DS-126c24f2-174c-46ca-b9bb-7ba343d48a2c,DISK], DatanodeInfoWithStorage[172.16.122.132:50010,DS-df4aac5a-b829-4b47-adbe-b5d79cbc75de,DISK], DatanodeInfoWithStorage[172.16.122.131:50010,DS-f4d957e2-1c71-42e9-8310-c5d0aec12cb7,DISK]]
1. BP-2089008039-172.16.122.134-1615302408196:blk_1073743599_2775 len=134217728 Live_repl=3 [DatanodeInfoWithStorage[172.16.122.132:50010,DS-4ef65136-5875-4096-a8df-b7d5644a133e,DISK], DatanodeInfoWithStorage[172.16.122.130:50010,DS-7440a047-980e-405f-aa45-d69fc6aa5396,DISK], DatanodeInfoWithStorage[172.16.122.131:50010,DS-d4df355f-be8b-4ca6-95b6-75b1a3549cb7,DISK]]
2. BP-2089008039-172.16.122.134-1615302408196:blk_1073743600_2776 len=134217728 Live_repl=3 [DatanodeInfoWithStorage[172.16.122.130:50010,DS-126c24f2-174c-46ca-b9bb-7ba343d48a2c,DISK], DatanodeInfoWithStorage[172.16.122.131:50010,DS-f4d957e2-1c71-42e9-8310-c5d0aec12cb7,DISK], DatanodeInfoWithStorage[172.16.122.132:50010,DS-df4aac5a-b829-4b47-adbe-b5d79cbc75de,DISK]]
3. BP-2089008039-172.16.122.134-1615302408196:blk_1073743601_2777 len=47547629 Live_repl=3 [DatanodeInfoWithStorage[172.16.122.130:50010,DS-7440a047-980e-405f-aa45-d69fc6aa5396,DISK], DatanodeInfoWithStorage[172.16.122.131:50010,DS-d4df355f-be8b-4ca6-95b6-75b1a3549cb7,DISK], DatanodeInfoWithStorage[172.16.122.132:50010,DS-4ef65136-5875-4096-a8df-b7d5644a133e,DISK]]
```

```
Status: HEALTHY
Number of data-nodes: 3
Number of racks: 1
Total dirs: 0
Total symlinks: 0
```

Replicated Blocks:

```
Total size: 450200813 B
Total files: 1
Total blocks (validated): 4 (avg. block size 112550203 B)
Minimally replicated blocks: 4 (100.0 %)
Over-replicated blocks: 0 (0.0 %)
Under-replicated blocks: 0 (0.0 %)
Mis-replicated blocks: 0 (0.0 %)
Default replication factor: 3
Average block replication: 3.0
Missing blocks: 0
Corrupt blocks: 0
Missing replicas: 0 (0.0 %)
```

→ $134,217,728 * 3 + 47,547,629 = 450,200,813$ bytes (450 MB)

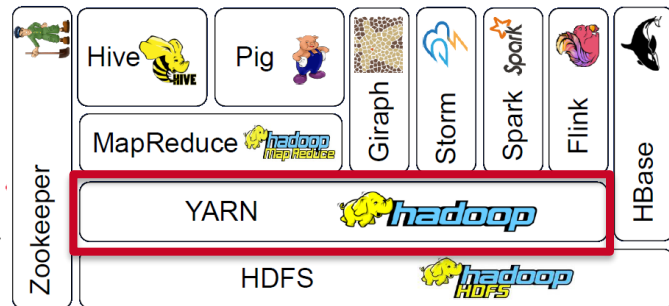
HDFS

- All the technicalities related to HDFS (configuration, load balancing, installation, *etc.*) are often handled by IT professionals
 - Data scientists use HDFS to request/send files (data) to/from Hadoop
- Analogy: accessing data/storing data from/in a relational database
 - Database administrator = IT professional who takes care of the technical aspects of running/maintaining the database

HDFS

- The Name Node can tell an application/task where each piece of data is stored
 - The relevant computations are then executed in the Data Node
 - “Moving computation to data”
 - Moving data around is costly (bandwidth wise)
- Unanswered questions
 - If a task wants to access the resources from a node (CPU, memory, disk, ...), how do we know whether that node is busy
 - Busy = dealing with another task
 - We need some sort of resource manager/negotiator!

YARN



➤ Yet Another Resource Negotiator

- Layer on top of HDFS
- Schedule resources to be used by different applications
 - Enables running multiple applications over HDFS in parallel

➤ In practice, it is rare to deal with YARN directly

- In this course, we will use Python to connect to high-level services (such as Spark), which in turn may use YARN behind the scenes

Hadoop

- When to use Hadoop HDFS
 - High data volume
 - High data variety
 - Multiple or parallel computations using the same data
- When not to use Hadoop
 - “Small data” processing
 - To store very well-structured data
 - Relational databases will likely do a better job
 - To store textual data only
 - MongoDB will likely do a better job

Summary

- We learned about a basic service in Hadoop
 - HDFS
 - Distributed file system
- Homework 10 is available on Canvas
- Next lecture
 - Spark